

Mini-project Group 12

January 2, 2024

Oskar Lundqvist; osauli-0@student.ltu.se

Filip Renberg; filren-0@student.ltu.se

We are gonna do a simple lab where we train a model using Keras

We are gonna use the mnist dataset of digits imported from keras like we did in lab 1 and lab 4.

```
[ ]: from silence_tensorflow import silence_tensorflow
silence_tensorflow() # Remove non-important tensorflow warnings

import numpy as np
import keras
from keras import layers
from keras.datasets import mnist
input_shape = (28,28,1)
num_classes = 10
```

<IPython.core.display.HTML object>

```
WARNING:tensorflow:From C:\Users\oskar\AppData\Local\Packages\PythonSoftwareFoun
dation.Python.3.11_qbz5n2kfra8p0\LocalCache\local-packages\Python311\site-
packages\keras\src\losses.py:2976: The name
tf.losses.sparse_softmax_cross_entropy is deprecated. Please use
tf.compat.v1.losses.sparse_softmax_cross_entropy instead.
```

We preprocess our data

```
[ ]: (train_data, train_labels), (test_data, test_labels) = mnist.load_data()
#Normalize input values
train_data = train_data.astype("float32")/255
test_data = test_data.astype("float32")/255
#Make sure images have the correct shape(28,28,1)
train_data = np.expand_dims(train_data, -1)
test_data = np.expand_dims(test_data, -1)

#Converts class vectors to binary class matrices
train_labels = keras.utils.to_categorical(train_labels, num_classes)
test_labels = keras.utils.to_categorical(test_labels, num_classes)
```

```

print("training shape: ", train_data.shape)
print(train_data.shape[0], "number of training samples")
print(test_data.shape[0], "number of testing samples")

```

training shape: (60000, 28, 28, 1)
 60000 number of training samples
 10000 number of testing samples

We create our convolutional neural network using Keras

```

[ ]: def create_model(nr_conv2d = 1, nr_filter = 32, kernel_size = (3, 3),
    ↪model_summary = False):
    # Name the model based on the input values,
    # Model_nr_conv2d-nr_filter-kernel_size
    model_name=f"sequential_{nr_conv2d}-{nr_filter}-{kernel_size[0]}"

    hidden_layers = []

    hidden_layers.append(keras.Input(shape=input_shape)) # input layer

    for i in range(nr_conv2d):
        hidden_layers.append(layers.Conv2D(nr_filter, kernel_size=kernel_size,
    ↪activation="relu"))
        hidden_layers.append(layers.MaxPooling2D(pool_size=[2,2]))

        nr_filter*=2

    hidden_layers.append(layers.Flatten())
    hidden_layers.append(layers.Dropout(0.5)) #Prevents overfitting
    hidden_layers.append(layers.Dense(num_classes, activation="softmax"))

    model = keras.Sequential(hidden_layers, name=model_name)

    if model_summary == True:
        model.summary()
    else:
        print(f'Model: {model_name}')

    return model

```

Next we are training the model. We can use several different loss functions here, for this project we will use “categorical_crossentropy”, “poisson” and “binary_crossentropy”

```

[ ]: n_epochs = 10
    size_batch = 128

    def model_training(pick, model):
        match pick:

```

```

    case 1:
        #Categorical loss function
        model.compile(loss="categorical_crossentropy", optimizer="adam",
↪metrics=['categorical_accuracy'])
        print("Loss Function: categorical_crossentropy")
        model.fit(train_data, train_labels, batch_size=size_batch,
↪epochs=n_epochs, validation_split=0.1, verbose=0)
    case 2:
        #Poisson loss function
        model.compile(loss="poisson", optimizer="adam",
↪metrics=['categorical_accuracy'])
        print("Loss Function: poisson")
        model.fit(train_data, train_labels, batch_size=size_batch,
↪epochs=n_epochs, validation_split=0.1, verbose=0)
    case 3:
        #Binary loss function
        model.compile(loss="binary_crossentropy", optimizer="adam",
↪metrics=['categorical_accuracy'])
        print("Loss Function: binary_crossentropy")
        model.fit(train_data, train_labels, batch_size=size_batch,
↪epochs=n_epochs, validation_split=0.1, verbose=0)

```

Define different values that we want to test

```

[ ]: nr_conv2d_list = [1, 2, 3]
    nr_filter_list = [8, 16, 32]
    kernel_size_list = [1, 2, 3]

    # We use this list to save the values for the best model
    best_model = [[0, 1]]

```

```

[ ]: def best_performance(best_model):
    # Resulting best model:
    create_model(nr_conv2d=best_model[0][2], nr_filter=best_model[0][3],
↪model_summary=True)
    match best_model[0][4]:
        case 1:
            print("Using Categorical crossentropy loss function resulted in the
↪\nfollowing accuracy and loss:")
        case 2:
            print("Using Poisson loss function resulted in the \nfollowing
↪accuracy and loss:")
        case 3:
            print("Using binary crossentropy loss function resulted in the
↪\nfollowing accuracy and loss")
            print("| \taccuracy: \t", best_model[0][0], "\t|")
            print("| \tloss: \t\t", best_model[0][1], "\t|")

```

```
[ ]: print("We mute model summary for the test prints but you can still see the_
      ↪model name:")
print("first value is the nr of conv2d layers, second value is the starting_
      ↪filter size.")
print("The third value is the kernel size but we don't change it in our current_
      ↪tests.\n")

for i in nr_conv2d_list: # The nr conv2d layers
    for j in nr_filter_list: # The starting size of the conv2d filters
        for k in kernel_size_list: # Kernel size for the filters
            for n in range(3): # The loss function

                model = create_model(nr_conv2d=i, nr_filter=j, kernel_size=(k,
            ↪k)) # You can also change the kernel size
                train = model_training(n+1, model)

                #Evaluate
                score = model.evaluate(test_data, test_labels, verbose=0)
                print("\taccuracy: \t", score[1], "\t|")
                print("\tloss: \t\t", score[0], "\t|")

                if (score[1] > best_model[0][0]) & (score[0] <
            ↪best_model[0][1]):
                    best_model.pop()
                    best_model.append((score[1], score[0], i, j, n+1))

                print("\n")

print("Finally done :)")
```

We mute model summary for the test prints but you can still see the model name:
first value is the nr of conv2d layers, second value is the starting filter
size.

The third value is the kernel size but we don't change it in our current tests.

Model: sequential_1-8-1

Loss Function: categorical_crossentropy

	accuracy:	0.9045000076293945	
	loss:	0.34677305817604065	

Model: sequential_1-8-1

Loss Function: poisson

	accuracy:	0.9099000096321106	
	loss:	0.13107535243034363	

Model: sequential_1-8-1
 Loss Function: binary_crossentropy

	accuracy:	0.9020000100135803	
	loss:	0.07336410880088806	

Model: sequential_1-8-2
 Loss Function: categorical_crossentropy

	accuracy:	0.958899974822998	
	loss:	0.14431360363960266	

Model: sequential_1-8-2
 Loss Function: poisson

	accuracy:	0.9645000100135803	
	loss:	0.11252903193235397	

Model: sequential_1-8-2
 Loss Function: binary_crossentropy

	accuracy:	0.9466999769210815	
	loss:	0.04450652003288269	

Model: sequential_1-8-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9711999893188477	
	loss:	0.09822211414575577	

Model: sequential_1-8-3
 Loss Function: poisson

	accuracy:	0.9735999703407288	
	loss:	0.10903801023960114	

Model: sequential_1-8-3
 Loss Function: binary_crossentropy

	accuracy:	0.965499997138977	
	loss:	0.02963523007929325	

Model: sequential_1-16-1
 Loss Function: categorical_crossentropy

	accuracy:	0.9160000085830688	
	loss:	0.2961479127407074	

Model: sequential_1-16-1
 Loss Function: poisson

	accuracy:	0.9136000275611877	
	loss:	0.12960928678512573	

Model: sequential_1-16-1
 Loss Function: binary_crossentropy

	accuracy:	0.9064000248908997	
	loss:	0.06932251155376434	

Model: sequential_1-16-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9725000262260437	
	loss:	0.09533493220806122	

Model: sequential_1-16-2
 Loss Function: poisson

	accuracy:	0.9718999862670898	
	loss:	0.10945045202970505	

Model: sequential_1-16-2
 Loss Function: binary_crossentropy

	accuracy:	0.96670001745224	
	loss:	0.028368938714265823	

Model: sequential_1-16-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9800000190734863	
	loss:	0.06785483658313751	

Model: sequential_1-16-3
 Loss Function: poisson

	accuracy:	0.9778000116348267	
	loss:	0.10673431307077408	

Model: sequential_1-16-3
 Loss Function: binary_crossentropy

	accuracy:	0.9771999716758728	
	loss:	0.02028021775186062	

Model: sequential_1-32-1
 Loss Function: categorical_crossentropy

	accuracy:	0.9169999957084656	
	loss:	0.2856709957122803	

Model: sequential_1-32-1
 Loss Function: poisson

	accuracy:	0.9178000092506409	
	loss:	0.1284821480512619	

Model: sequential_1-32-1
 Loss Function: binary_crossentropy

	accuracy:	0.9103000164031982	
	loss:	0.06726555526256561	

Model: sequential_1-32-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9753000140190125	
	loss:	0.08092426508665085	

Model: sequential_1-32-2
 Loss Function: poisson

	accuracy:	0.9750999808311462	
	loss:	0.10779019445180893	

Model: sequential_1-32-2
 Loss Function: binary_crossentropy

	accuracy:	0.9740999937057495	
	loss:	0.02248363569378853	

Model: sequential_1-32-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9822999835014343	
	loss:	0.05565842241048813	

Model: sequential_1-32-3
 Loss Function: poisson

	accuracy:	0.9824000000953674	
	loss:	0.10548123717308044	

Model: sequential_1-32-3
 Loss Function: binary_crossentropy

	accuracy:	0.9793000221252441	
	loss:	0.01774715445935726	

Model: sequential_2-8-1
 Loss Function: categorical_crossentropy

	accuracy:	0.862500011920929	
	loss:	0.4641565978527069	

Model: sequential_2-8-1
 Loss Function: poisson

	accuracy:	0.8756999969482422	
	loss:	0.1407691091299057	

Model: sequential_2-8-1
 Loss Function: binary_crossentropy

	accuracy:	0.8684999942779541	
	loss:	0.09200688451528549	

Model: sequential_2-8-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9779999852180481	
	loss:	0.06890183687210083	

Model: sequential_2-8-2
 Loss Function: poisson

	accuracy:	0.9779000282287598	
	loss:	0.10701524466276169	

Model: sequential_2-8-2
 Loss Function: binary_crossentropy

	accuracy:	0.9751999974250793	
	loss:	0.023338187485933304	

Model: sequential_2-8-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9829000234603882	
	loss:	0.055518683046102524	

Model: sequential_2-8-3
 Loss Function: poisson
 | accuracy: 0.984000027179718 |
 | loss: 0.1050797626376152 |

Model: sequential_2-8-3
 Loss Function: binary_crossentropy
 | accuracy: 0.9811000227928162 |
 | loss: 0.016874121502041817 |

Model: sequential_2-16-1
 Loss Function: categorical_crossentropy
 | accuracy: 0.8657000064849854 |
 | loss: 0.43340224027633667 |

Model: sequential_2-16-1
 Loss Function: poisson
 | accuracy: 0.8867999911308289 |
 | loss: 0.1376599669456482 |

Model: sequential_2-16-1
 Loss Function: binary_crossentropy
 | accuracy: 0.8590999841690063 |
 | loss: 0.09501325339078903 |

Model: sequential_2-16-2
 Loss Function: categorical_crossentropy
 | accuracy: 0.9840999841690063 |
 | loss: 0.05054925009608269 |

Model: sequential_2-16-2
 Loss Function: poisson
 | accuracy: 0.9824000000953674 |
 | loss: 0.10543941706418991 |

Model: sequential_2-16-2
 Loss Function: binary_crossentropy
 | accuracy: 0.9825999736785889 |
 | loss: 0.015554818324744701 |

Model: sequential_2-16-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9876000285148621	
	loss:	0.035443928092718124	

Model: sequential_2-16-3
 Loss Function: poisson

	accuracy:	0.9896000027656555	
	loss:	0.10363186150789261	

Model: sequential_2-16-3
 Loss Function: binary_crossentropy

	accuracy:	0.9871000051498413	
	loss:	0.010172258131206036	

Model: sequential_2-32-1
 Loss Function: categorical_crossentropy

	accuracy:	0.8823000192642212	
	loss:	0.38418707251548767	

Model: sequential_2-32-1
 Loss Function: poisson

	accuracy:	0.8726999759674072	
	loss:	0.14082711935043335	

Model: sequential_2-32-1
 Loss Function: binary_crossentropy

	accuracy:	0.8611999750137329	
	loss:	0.09058661758899689	

Model: sequential_2-32-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9868000149726868	
	loss:	0.040284544229507446	

Model: sequential_2-32-2
 Loss Function: poisson

	accuracy:	0.9854999780654907	
	loss:	0.1041753813624382	

Model: sequential_2-32-2
 Loss Function: binary_crossentropy

	accuracy:	0.9861999750137329	
	loss:	0.01182855386286974	

Model: sequential_2-32-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9902999997138977	
	loss:	0.029759705066680908	

Model: sequential_2-32-3
 Loss Function: poisson

	accuracy:	0.9919000267982483	
	loss:	0.10243698954582214	

Model: sequential_2-32-3
 Loss Function: binary_crossentropy

	accuracy:	0.989799976348877	
	loss:	0.008357501588761806	

Model: sequential_3-8-1
 Loss Function: categorical_crossentropy

	accuracy:	0.6349999904632568	
	loss:	1.1341913938522339	

Model: sequential_3-8-1
 Loss Function: poisson

	accuracy:	0.6380000114440918	
	loss:	0.21298415958881378	

Model: sequential_3-8-1
 Loss Function: binary_crossentropy

	accuracy:	0.6355000138282776	
	loss:	0.19037678837776184	

Model: sequential_3-8-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9696000218391418	
	loss:	0.10524994134902954	

Model: sequential_3-8-2
 Loss Function: poisson

	accuracy:	0.9692999720573425	
	loss:	0.11103629320859909	

Model: sequential_3-8-2
 Loss Function: binary_crossentropy

	accuracy:	0.9581999778747559	
	loss:	0.035895027220249176	

Model: sequential_3-8-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9672999978065491	
	loss:	0.11193626374006271	

Model: sequential_3-8-3
 Loss Function: poisson

	accuracy:	0.9697999954223633	
	loss:	0.11056715250015259	

Model: sequential_3-8-3
 Loss Function: binary_crossentropy

	accuracy:	0.9596999883651733	
	loss:	0.038717880845069885	

Model: sequential_3-16-1
 Loss Function: categorical_crossentropy

	accuracy:	0.661899983882904	
	loss:	1.0449005365371704	

Model: sequential_3-16-1
 Loss Function: poisson

	accuracy:	0.6635000109672546	
	loss:	0.20406877994537354	

Model: sequential_3-16-1
 Loss Function: binary_crossentropy

	accuracy:	0.6579999923706055	
	loss:	0.18362551927566528	

Model: sequential_3-16-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9815999865531921	
	loss:	0.05796007812023163	

Model: sequential_3-16-2
 Loss Function: poisson

	accuracy:	0.9830999970436096	
	loss:	0.1051589846611023	

Model: sequential_3-16-2
 Loss Function: binary_crossentropy

	accuracy:	0.9789000153541565	
	loss:	0.01722176931798458	

Model: sequential_3-16-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9839000105857849	
	loss:	0.05719917640089989	

Model: sequential_3-16-3
 Loss Function: poisson

	accuracy:	0.984000027179718	
	loss:	0.1055717021226883	

Model: sequential_3-16-3
 Loss Function: binary_crossentropy

	accuracy:	0.9797999858856201	
	loss:	0.016070352867245674	

Model: sequential_3-32-1
 Loss Function: categorical_crossentropy

	accuracy:	0.6797999739646912	
	loss:	0.9819020628929138	

Model: sequential_3-32-1
 Loss Function: poisson

	accuracy:	0.6796000003814697	
	loss:	0.19981011748313904	

```

Model: sequential_3-32-1
Loss Function: binary_crossentropy
|      accuracy:      0.6588000059127808      |
|      loss:          0.17940448224544525      |

```

```

Model: sequential_3-32-2
Loss Function: categorical_crossentropy
|      accuracy:      0.9872999787330627      |
|      loss:          0.04084913805127144      |

```

```

Model: sequential_3-32-2
Loss Function: poisson
|      accuracy:      0.9879000186920166      |
|      loss:          0.10404383391141891      |

```

```

Model: sequential_3-32-2
Loss Function: binary_crossentropy
|      accuracy:      0.9864000082015991      |
|      loss:          0.011171533726155758      |

```

```

Model: sequential_3-32-3
Loss Function: categorical_crossentropy
|      accuracy:      0.9884999990463257      |
|      loss:          0.03939943015575409      |

```

```

Model: sequential_3-32-3
Loss Function: poisson
|      accuracy:      0.9868000149726868      |
|      loss:          0.10447220504283905      |

```

```

Model: sequential_3-32-3
Loss Function: binary_crossentropy
|      accuracy:      0.9861999750137329      |
|      loss:          0.010479954071342945      |

```

Finally done :)

The following is the best performing model from our tests:

```
[ ]: best_performance(best_model)
```

```
Model: "sequential_2-32-3"
```

Layer (type)	Output Shape	Param #
conv2d_162 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_162 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_163 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_163 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_81 (Flatten)	(None, 1600)	0
dropout_81 (Dropout)	(None, 1600)	0
dense_81 (Dense)	(None, 10)	16010

Total params: 34826 (136.04 KB)
 Trainable params: 34826 (136.04 KB)
 Non-trainable params: 0 (0.00 Byte)

Using Categorical crossentropy loss function resulted in the following accuracy and loss:

	accuracy:	0.9902999997138977	
	loss:	0.029759705066680908	

Layer (type)	Output Shape	Param #
conv2d_162 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_162 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_163 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_163 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_81 (Flatten)	(None, 1600)	0
dropout_81 (Dropout)	(None, 1600)	0
dense_81 (Dense)	(None, 10)	16010

Total params: 34826 (136.04 KB)
Trainable params: 34826 (136.04 KB)
Non-trainable params: 0 (0.00 Byte)

Using Categorical crossentropy loss function resulted in the
following accuracy and loss:

	accuracy:	0.9902999997138977	
	loss:	0.029759705066680908	