

Mini-project Group 12

January 6, 2024

Oskar Lundqvist; osauli-0@student.ltu.se

Filip Renberg; filren-0@student.ltu.se

We are gonna do a simple lab where we train a model using Keras

We are gonna use the mnist dataset of digits imported from keras like we did in lab 1 and lab 4.

```
[ ]: import logging
logging.getLogger('tensorflow').disabled = True # Disable warnings from
↳ tensorflow
import numpy as np
import keras
from keras import layers
from keras.datasets import mnist
input_shape = (28,28,1)
num_classes = 10
```

We preprocess our data

```
[ ]: (train_data, train_labels), (test_data, test_labels) = mnist.load_data()
#Normalize input values
train_data = train_data.astype("float32")/255
test_data = test_data.astype("float32")/255
#Make sure images have the correct shape(28,28,1)
train_data = np.expand_dims(train_data, -1)
test_data = np.expand_dims(test_data, -1)

#Converts class vectors to binary class matrices
train_labels = keras.utils.to_categorical(train_labels, num_classes)
test_labels = keras.utils.to_categorical(test_labels, num_classes)

print("training shape: ", train_data.shape)
print(train_data.shape[0], "number of training samples")
print(test_data.shape[0], "number of testing samples")
```

training shape: (60000, 28, 28, 1)

60000 number of training samples

10000 number of testing samples

We create our convolutional neural network using Keras

```
[ ]: def create_model(nr_conv2d = 1, nr_filter = 32, kernel_size = (3, 3),
    ↪model_summary = False):
    # Name the model based on the input values,
    # Sequential_nr_conv2d-nr_filter-kernel_size
    model_name=f"sequential_{nr_conv2d}-{nr_filter}-{kernel_size[0]}"

    hidden_layers = []

    hidden_layers.append(keras.Input(shape=input_shape)) # input layer

    for i in range(nr_conv2d):
        hidden_layers.append(layers.Conv2D(nr_filter, kernel_size=kernel_size,
    ↪activation="relu"))
        hidden_layers.append(layers.MaxPooling2D(pool_size=[2,2]))

        nr_filter*=2

    hidden_layers.append(layers.Flatten())
    hidden_layers.append(layers.Dropout(0.5)) #Prevents overfitting
    hidden_layers.append(layers.Dense(num_classes, activation="softmax"))

    model = keras.Sequential(hidden_layers, name=model_name)

    if model_summary == True:
        model.summary()
    else:
        print(f'Model: {model_name}')

    return model
```

Next we are training the model. We can use several different loss functions here, for this project we will use “categorical_crossentropy”, “poisson” and “binary_crossentropy”

```
[ ]: n_epochs = 10
    size_batch = 128

    def model_training(pick, model):
        match pick:
            case 1:
                #Categorical loss function
                model.compile(loss="categorical_crossentropy", optimizer="adam",
    ↪metrics=['categorical_accuracy'])
                print("Loss Function: categorical_crossentropy")
                model.fit(train_data, train_labels, batch_size=size_batch,
    ↪epochs=n_epochs, validation_split=0.1, verbose=0)
            case 2:
                #Poisson loss function
```

```

        model.compile(loss="poisson", optimizer="adam",
↳metrics=['categorical_accuracy'])
        print("Loss Function: poisson")
        model.fit(train_data, train_labels, batch_size=size_batch,
↳epochs=n_epochs, validation_split=0.1, verbose=0)
        case 3:
            #Binary loss function
            model.compile(loss="binary_crossentropy", optimizer="adam",
↳metrics=['categorical_accuracy'])
            print("Loss Function: binary_crossentropy")
            model.fit(train_data, train_labels, batch_size=size_batch,
↳epochs=n_epochs, validation_split=0.1, verbose=0)

```

Define different values that we want to test

```

[ ]: nr_conv2d_list = [1, 2, 3]
     nr_filter_list = [8, 16, 32]
     kernel_size_list = [1, 2, 3]

     # We use this list to save the values for the best performing models
     best_overall = [[0, 1]]
     best_accuracy = [[0, 1]]
     best_loss = [[0, 1]]

```

```

[ ]: def best_performance(best_model):
     # Resulting best model:
     create_model(nr_conv2d=best_model[0][2], nr_filter=best_model[0][3],
↳model_summary=True)
     match best_model[0][4]:
         case 1:
             print("Using Categorical crossentropy loss function resulted in the
↳\nfollowing accuracy and loss:")
         case 2:
             print("Using Poisson loss function resulted in the \nfollowing
↳accuracy and loss:")
         case 3:
             print("Using binary crossentropy loss function resulted in the
↳\nfollowing accuracy and loss")
             print("| \taccuracy: \t", best_model[0][0], "\t|")
             print("| \tloss: \t\t", best_model[0][1], "\t|")

```

We mute model summary for the test prints but you can still see the model name, example:

sequential_x1-x2-x3

x1 is the nr of conv2d layers, x2 is the starting filter size, and x3 is the kernel size.

```
[ ]: for i in nr_conv2d_list: # The nr conv2d layers
      for j in nr_filter_list: # The starting size of the conv2d filters
          for k in kernel_size_list: # Kernel size for the filters
              for n in range(3): # The loss function

                  model = create_model(nr_conv2d=i, nr_filter=j, kernel_size=(k,
↪k)) # You can also change the kernel size
                  train = model_training(n+1, model)

                  #Evaluate
                  score = model.evaluate(test_data, test_labels, verbose=0)
                  print("\taccuracy: \t", score[1], "\t|")
                  print("\tloss: \t\t", score[0], "\t|")

                  if (score[1] > best_overall[0][0]) & (score[0] <
↪best_overall[0][1]): # best accuracy + loss
                      best_overall.pop()
                      best_overall.append((score[1], score[0], i, j, n+1))

                  if (score[1] > best_accuracy[0][0]): # best accuracy
                      best_accuracy.pop()
                      best_accuracy.append((score[1], score[0], i, j, n+1))

                  if (score[0] < best_loss[0][1]): # best loss
                      best_loss.pop()
                      best_loss.append((score[1], score[0], i, j, n+1))

                  print("\n")

print("Finally done :)")
```

Model: sequential_1-8-1

Loss Function: categorical_crossentropy

	accuracy:	0.9111999869346619	
	loss:	0.30452200770378113	

Model: sequential_1-8-1

Loss Function: poisson

	accuracy:	0.9142000079154968	
	loss:	0.13022562861442566	

Model: sequential_1-8-1

Loss Function: binary_crossentropy

	accuracy:	0.8944000005722046	
	loss:	0.08388058841228485	

Model: sequential_1-8-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9581999778747559	
	loss:	0.14477361738681793	

Model: sequential_1-8-2
 Loss Function: poisson

	accuracy:	0.961899995803833	
	loss:	0.11375628411769867	

Model: sequential_1-8-2
 Loss Function: binary_crossentropy

	accuracy:	0.954200029373169	
	loss:	0.039917927235364914	

Model: sequential_1-8-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9682999849319458	
	loss:	0.10456815361976624	

Model: sequential_1-8-3
 Loss Function: poisson

	accuracy:	0.9692999720573425	
	loss:	0.11055201292037964	

Model: sequential_1-8-3
 Loss Function: binary_crossentropy

	accuracy:	0.9617999792098999	
	loss:	0.031008990481495857	

Model: sequential_1-16-1
 Loss Function: categorical_crossentropy

	accuracy:	0.9124000072479248	
	loss:	0.2945832312107086	

Model: sequential_1-16-1
 Loss Function: poisson

	accuracy:	0.914900004863739	
	loss:	0.1297198385000229	

Model: sequential_1-16-1
 Loss Function: binary_crossentropy

	accuracy:	0.9083999991416931	
	loss:	0.06987139582633972	

Model: sequential_1-16-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9692999720573425	
	loss:	0.10458052903413773	

Model: sequential_1-16-2
 Loss Function: poisson

	accuracy:	0.9718999862670898	
	loss:	0.10959072411060333	

Model: sequential_1-16-2
 Loss Function: binary_crossentropy

	accuracy:	0.9621000289916992	
	loss:	0.03141861408948898	

Model: sequential_1-16-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9761999845504761	
	loss:	0.07268141210079193	

Model: sequential_1-16-3
 Loss Function: poisson

	accuracy:	0.9786999821662903	
	loss:	0.10692232847213745	

Model: sequential_1-16-3
 Loss Function: binary_crossentropy

	accuracy:	0.9765999913215637	
	loss:	0.020543649792671204	

Model: sequential_1-32-1
 Loss Function: categorical_crossentropy

	accuracy:	0.9176999926567078	
	loss:	0.28041988611221313	

Model: sequential_1-32-1
 Loss Function: poisson

	accuracy:	0.9190000295639038	
	loss:	0.1282191425561905	

Model: sequential_1-32-1
 Loss Function: binary_crossentropy

	accuracy:	0.9085000157356262	
	loss:	0.06618037074804306	

Model: sequential_1-32-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9760000109672546	
	loss:	0.07642963528633118	

Model: sequential_1-32-2
 Loss Function: poisson

	accuracy:	0.9764000177383423	
	loss:	0.10754676163196564	

Model: sequential_1-32-2
 Loss Function: binary_crossentropy

	accuracy:	0.9682000279426575	
	loss:	0.025774776935577393	

Model: sequential_1-32-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9824000000953674	
	loss:	0.056834980845451355	

Model: sequential_1-32-3
 Loss Function: poisson

	accuracy:	0.9814000129699707	
	loss:	0.1056058332324028	

Model: sequential_1-32-3
 Loss Function: binary_crossentropy

	accuracy:	0.9781000018119812	
	loss:	0.01797471195459366	

Model: sequential_2-8-1
 Loss Function: categorical_crossentropy
 | accuracy: 0.8546000123023987 |
 | loss: 0.4925113916397095 |

Model: sequential_2-8-1
 Loss Function: poisson
 | accuracy: 0.8597000241279602 |
 | loss: 0.14700977504253387 |

Model: sequential_2-8-1
 Loss Function: binary_crossentropy
 | accuracy: 0.8456000089645386 |
 | loss: 0.1057511568069458 |

Model: sequential_2-8-2
 Loss Function: categorical_crossentropy
 | accuracy: 0.975600004196167 |
 | loss: 0.0781906396150589 |

Model: sequential_2-8-2
 Loss Function: poisson
 | accuracy: 0.9761999845504761 |
 | loss: 0.10781963169574738 |

Model: sequential_2-8-2
 Loss Function: binary_crossentropy
 | accuracy: 0.9746999740600586 |
 | loss: 0.023071149364113808 |

Model: sequential_2-8-3
 Loss Function: categorical_crossentropy
 | accuracy: 0.9829999804496765 |
 | loss: 0.0523674450814724 |

Model: sequential_2-8-3
 Loss Function: poisson
 | accuracy: 0.9837999939918518 |
 | loss: 0.1052933856844902 |

Model: sequential_2-8-3
 Loss Function: binary_crossentropy

	accuracy:	0.982699990272522	
	loss:	0.01610827073454857	

Model: sequential_2-16-1
 Loss Function: categorical_crossentropy

	accuracy:	0.8632000088691711	
	loss:	0.45609068870544434	

Model: sequential_2-16-1
 Loss Function: poisson

	accuracy:	0.8676999807357788	
	loss:	0.14288942515850067	

Model: sequential_2-16-1
 Loss Function: binary_crossentropy

	accuracy:	0.8587999939918518	
	loss:	0.09436600655317307	

Model: sequential_2-16-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9837999939918518	
	loss:	0.048797789961099625	

Model: sequential_2-16-2
 Loss Function: poisson

	accuracy:	0.9842000007629395	
	loss:	0.10497941821813583	

Model: sequential_2-16-2
 Loss Function: binary_crossentropy

	accuracy:	0.9804999828338623	
	loss:	0.01698094978928566	

Model: sequential_2-16-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9890000224113464	
	loss:	0.03584475442767143	

Model: sequential_2-16-3
 Loss Function: poisson

	accuracy:	0.9879000186920166	
	loss:	0.10372993350028992	

Model: sequential_2-16-3
 Loss Function: binary_crossentropy

	accuracy:	0.987500011920929	
	loss:	0.011103938333690166	

Model: sequential_2-32-1
 Loss Function: categorical_crossentropy

	accuracy:	0.8661999702453613	
	loss:	0.42086517810821533	

Model: sequential_2-32-1
 Loss Function: poisson

	accuracy:	0.8622999787330627	
	loss:	0.1432160884141922	

Model: sequential_2-32-1
 Loss Function: binary_crossentropy

	accuracy:	0.866100013256073	
	loss:	0.08706800639629364	

Model: sequential_2-32-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9861000180244446	
	loss:	0.04110102355480194	

Model: sequential_2-32-2
 Loss Function: poisson

	accuracy:	0.9872000217437744	
	loss:	0.10391152650117874	

Model: sequential_2-32-2
 Loss Function: binary_crossentropy

	accuracy:	0.98580002784729	
	loss:	0.011923989281058311	

Model: sequential_2-32-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9912999868392944	
	loss:	0.025549551472067833	

Model: sequential_2-32-3
 Loss Function: poisson

	accuracy:	0.9907000064849854	
	loss:	0.10306503623723984	

Model: sequential_2-32-3
 Loss Function: binary_crossentropy

	accuracy:	0.9904999732971191	
	loss:	0.008173882029950619	

Model: sequential_3-8-1
 Loss Function: categorical_crossentropy

	accuracy:	0.566100001335144	
	loss:	1.3437367677688599	

Model: sequential_3-8-1
 Loss Function: poisson

	accuracy:	0.6552000045776367	
	loss:	0.20856887102127075	

Model: sequential_3-8-1
 Loss Function: binary_crossentropy

	accuracy:	0.633899986743927	
	loss:	0.19226402044296265	

Model: sequential_3-8-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9714000225067139	
	loss:	0.094182088971138	

Model: sequential_3-8-2
 Loss Function: poisson

	accuracy:	0.9739999771118164	
	loss:	0.10907341539859772	

Model: sequential_3-8-2
 Loss Function: binary_crossentropy

	accuracy:	0.9641000032424927	
	loss:	0.029729217290878296	

Model: sequential_3-8-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9692999720573425	
	loss:	0.09702785313129425	

Model: sequential_3-8-3
 Loss Function: poisson

	accuracy:	0.968500018119812	
	loss:	0.11046236753463745	

Model: sequential_3-8-3
 Loss Function: binary_crossentropy

	accuracy:	0.9638000130653381	
	loss:	0.03126277029514313	

Model: sequential_3-16-1
 Loss Function: categorical_crossentropy

	accuracy:	0.6657000184059143	
	loss:	1.0308948755264282	

Model: sequential_3-16-1
 Loss Function: poisson

	accuracy:	0.6869999766349792	
	loss:	0.19946211576461792	

Model: sequential_3-16-1
 Loss Function: binary_crossentropy

	accuracy:	0.6517000198364258	
	loss:	0.18309813737869263	

Model: sequential_3-16-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9818000197410583	
	loss:	0.06009068712592125	

Model: sequential_3-16-2
 Loss Function: poisson

	accuracy:	0.9811000227928162	
	loss:	0.10630549490451813	

Model: sequential_3-16-2
 Loss Function: binary_crossentropy

	accuracy:	0.980400025844574	
	loss:	0.01608153246343136	

Model: sequential_3-16-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9818999767303467	
	loss:	0.061925407499074936	

Model: sequential_3-16-3
 Loss Function: poisson

	accuracy:	0.9825999736785889	
	loss:	0.10590063780546188	

Model: sequential_3-16-3
 Loss Function: binary_crossentropy

	accuracy:	0.9768000245094299	
	loss:	0.017312271520495415	

Model: sequential_3-32-1
 Loss Function: categorical_crossentropy

	accuracy:	0.7143999934196472	
	loss:	0.9123744368553162	

Model: sequential_3-32-1
 Loss Function: poisson

	accuracy:	0.6766999959945679	
	loss:	0.19979707896709442	

Model: sequential_3-32-1
 Loss Function: binary_crossentropy

	accuracy:	0.6539999842643738	
	loss:	0.18011701107025146	

```

Model: sequential_3-32-2
Loss Function: categorical_crossentropy
|      accuracy:      0.9887999892234802      |
|      loss:          0.037356436252593994      |

```

```

Model: sequential_3-32-2
Loss Function: poisson
|      accuracy:      0.9861999750137329      |
|      loss:          0.10472308099269867      |

```

```

Model: sequential_3-32-2
Loss Function: binary_crossentropy
|      accuracy:      0.9864000082015991      |
|      loss:          0.010712560266256332      |

```

```

Model: sequential_3-32-3
Loss Function: categorical_crossentropy
|      accuracy:      0.9894999861717224      |
|      loss:          0.037665847688913345      |

```

```

Model: sequential_3-32-3
Loss Function: poisson
|      accuracy:      0.9871000051498413      |
|      loss:          0.10423722118139267      |

```

```

Model: sequential_3-32-3
Loss Function: binary_crossentropy
|      accuracy:      0.9857000112533569      |
|      loss:          0.010064337402582169      |

```

Finally done :)

The best overall performing model:

```
[ ]: best_performance(best_overall)
```

```
Model: "sequential_2-32-3"
```

Layer (type)	Output Shape	Param #
conv2d_162 (Conv2D)	(None, 26, 26, 32)	320

max_pooling2d_162 (MaxPool ing2D)	(None, 13, 13, 32)	0
conv2d_163 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_163 (MaxPool ing2D)	(None, 5, 5, 64)	0
flatten_81 (Flatten)	(None, 1600)	0
dropout_81 (Dropout)	(None, 1600)	0
dense_81 (Dense)	(None, 10)	16010

=====

Total params: 34826 (136.04 KB)
Trainable params: 34826 (136.04 KB)
Non-trainable params: 0 (0.00 Byte)

Using binary crossentropy loss function resulted in the
following accuracy and loss

accuracy:	0.9904999732971191	
loss:	0.008173882029950619	

Layer (type)	Output Shape	Param #
conv2d_162 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_162 (MaxPool ing2D)	(None, 13, 13, 32)	0
conv2d_163 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_163 (MaxPool ing2D)	(None, 5, 5, 64)	0
flatten_81 (Flatten)	(None, 1600)	0
dropout_81 (Dropout)	(None, 1600)	0
dense_81 (Dense)	(None, 10)	16010

=====

Total params: 34826 (136.04 KB)
Trainable params: 34826 (136.04 KB)
Non-trainable params: 0 (0.00 Byte)

Using binary crossentropy loss function resulted in the

following accuracy and loss

```
|      accuracy:      0.9904999732971191      |
|      loss:          0.008173882029950619      |
```

Most accurate model:

```
[ ]: best_performance(best_accuracy)
```

Model: "sequential_2-32-3"

Layer (type)	Output Shape	Param #
conv2d_164 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_164 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_165 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_165 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_82 (Flatten)	(None, 1600)	0
dropout_82 (Dropout)	(None, 1600)	0
dense_82 (Dense)	(None, 10)	16010

Total params: 34826 (136.04 KB)

Trainable params: 34826 (136.04 KB)

Non-trainable params: 0 (0.00 Byte)

Using Categorical crossentropy loss function resulted in the following accuracy and loss:

```
|      accuracy:      0.9912999868392944      |
|      loss:          0.025549551472067833      |
```

Layer (type)	Output Shape	Param #
conv2d_164 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_164 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_165 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_165 (MaxPooling2D)	(None, 5, 5, 64)	0

flatten_82 (Flatten)	(None, 1600)	0
dropout_82 (Dropout)	(None, 1600)	0
dense_82 (Dense)	(None, 10)	16010

```
=====
Total params: 34826 (136.04 KB)
Trainable params: 34826 (136.04 KB)
Non-trainable params: 0 (0.00 Byte)
```

```
-----
Using Categorical crossentropy loss function resulted in the
following accuracy and loss:
|      accuracy:      0.9912999868392944      |
|      loss:          0.025549551472067833      |
```

Best loss model:

```
[ ]: best_performance(best_loss)
```

Model: "sequential_2-32-3"

Layer (type)	Output Shape	Param #
conv2d_166 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_166 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_167 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_167 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_83 (Flatten)	(None, 1600)	0
dropout_83 (Dropout)	(None, 1600)	0
dense_83 (Dense)	(None, 10)	16010

```
=====
Total params: 34826 (136.04 KB)
Trainable params: 34826 (136.04 KB)
Non-trainable params: 0 (0.00 Byte)
```

```
-----
Using binary crossentropy loss function resulted in the
following accuracy and loss
|      accuracy:      0.9904999732971191      |
```

| loss: 0.008173882029950619 |

Layer (type)	Output Shape	Param #
conv2d_166 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_166 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_167 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_167 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_83 (Flatten)	(None, 1600)	0
dropout_83 (Dropout)	(None, 1600)	0
dense_83 (Dense)	(None, 10)	16010

=====
Total params: 34826 (136.04 KB)
Trainable params: 34826 (136.04 KB)
Non-trainable params: 0 (0.00 Byte)

Using binary_crossentropy loss function resulted in the
following accuracy and loss

| accuracy: 0.9904999732971191 |
| loss: 0.008173882029950619 |