

Mini-project Group 12

January 6, 2024

Oskar Lundqvist; osauli-0@student.ltu.se

Filip Renberg; filren-0@student.ltu.se

We are gonna do a simple lab where we train a model using Keras

We are gonna use the mnist dataset of digits imported from keras like we did in lab 1 and lab 4.

```
[ ]: import logging
logging.getLogger('tensorflow').disabled = True # Disable warnings from
↳ tensorflow
import numpy as np
import keras
from keras import layers
from keras.datasets import mnist
input_shape = (28,28,1)
num_classes = 10
```

We preprocess our data

```
[ ]: (train_data, train_labels), (test_data, test_labels) = mnist.load_data()
#Normalize input values
train_data = train_data.astype("float32")/255
test_data = test_data.astype("float32")/255
#Make sure images have the correct shape(28,28,1)
train_data = np.expand_dims(train_data, -1)
test_data = np.expand_dims(test_data, -1)

#Converts class vectors to binary class matrices
train_labels = keras.utils.to_categorical(train_labels, num_classes)
test_labels = keras.utils.to_categorical(test_labels, num_classes)

print("training shape: ", train_data.shape)
print(train_data.shape[0], "number of training samples")
print(test_data.shape[0], "number of testing samples")
```

training shape: (60000, 28, 28, 1)

60000 number of training samples

10000 number of testing samples

We create our convolutional neural network using Keras

```
[ ]: def create_model(nr_conv2d = 1, nr_filter = 32, kernel_size = (3, 3),
    ↪model_summary = False):
    # Name the model based on the input values,
    # Sequential_nr_conv2d-nr_filter-kernel_size
    model_name=f"sequential_{nr_conv2d}-{nr_filter}-{kernel_size[0]}"

    hidden_layers = []

    hidden_layers.append(keras.Input(shape=input_shape)) # input layer

    for i in range(nr_conv2d):
        hidden_layers.append(layers.Conv2D(nr_filter, kernel_size=kernel_size,
    ↪activation="relu"))
        hidden_layers.append(layers.MaxPooling2D(pool_size=[2,2]))

        nr_filter*=2

    hidden_layers.append(layers.Flatten())
    hidden_layers.append(layers.Dropout(0.5)) #Prevents overfitting
    hidden_layers.append(layers.Dense(num_classes, activation="softmax"))

    model = keras.Sequential(hidden_layers, name=model_name)

    if model_summary == True:
        model.summary()
    else:
        print(f'Model: {model_name}')

    return model
```

Next we are training the model. We can use several different loss functions here, for this project we will use “categorical_crossentropy”, “poisson” and “binary_crossentropy”

```
[ ]: n_epochs = 10
    size_batch = 128

    def model_training(pick, model):
        match pick:
            case 1:
                #Categorical loss function
                model.compile(loss="categorical_crossentropy", optimizer="adam",
    ↪metrics=['categorical_accuracy'])
                print("Loss Function: categorical_crossentropy")
                model.fit(train_data, train_labels, batch_size=size_batch,
    ↪epochs=n_epochs, validation_split=0.1, verbose=0)
            case 2:
                #Poisson loss function
```

```

        model.compile(loss="poisson", optimizer="adam",
↪metrics=['categorical_accuracy'])
        print("Loss Function: poisson")
        model.fit(train_data, train_labels, batch_size=size_batch,
↪epochs=n_epochs, validation_split=0.1, verbose=0)
        case 3:
            #Binary loss function
            model.compile(loss="binary_crossentropy", optimizer="adam",
↪metrics=['categorical_accuracy'])
            print("Loss Function: binary_crossentropy")
            model.fit(train_data, train_labels, batch_size=size_batch,
↪epochs=n_epochs, validation_split=0.1, verbose=0)

```

Define different values that we want to test

```

[ ]: nr_conv2d_list = [1, 2, 3]
     nr_filter_list = [8, 16, 32]
     kernel_size_list = [1, 2, 3]

     # We use this list to save the values for the best model
     best_model = [[0, 1]]

```

```

[ ]: def best_performance(best_model):
     # Resulting best model:
     create_model(nr_conv2d=best_model[0][2], nr_filter=best_model[0][3],
↪model_summary=True)
     match best_model[0][4]:
         case 1:
             print("Using Categorical crossentropy loss function resulted in the
↪\nfollowing accuracy and loss:")
         case 2:
             print("Using Poisson loss function resulted in the \nfollowing
↪accuracy and loss:")
         case 3:
             print("Using binary crossentropy loss function resulted in the
↪\nfollowing accuracy and loss")
             print("\t\taccuracy: \t", best_model[0][0], "\t\t")
             print("\t\tloss: \t\t", best_model[0][1], "\t\t")

```

We mute model summary for the test prints but you can still see the model name, example:

sequential_x1-x2-x3

x1 is the nr of conv2d layers, x2 is the starting filter size, and x3 is the kernel size.

```

[ ]: for i in nr_conv2d_list: # The nr conv2d layers
     for j in nr_filter_list: # The starting size of the conv2d filters
         for k in kernel_size_list: # Kernel size for the filters
             for n in range(3): # The loss function

```

```

        model = create_model(nr_conv2d=i, nr_filter=j, kernel_size=(k,
↪k)) # You can also change the kernel size
        train = model_training(n+1, model)

        #Evaluate
        score = model.evaluate(test_data, test_labels, verbose=0)
        print("\taccuracy: \t", score[1], "\t|")
        print("\tloss: \t\t", score[0], "\t|")

        if (score[1] > best_model[0][0]) & (score[0] <
↪best_model[0][1]):
            best_model.pop()
            best_model.append((score[1], score[0], i, j, n+1))

        print("\n")

print("Finally done :)")

```

Model: sequential_1-8-1

Loss Function: categorical_crossentropy

	accuracy:	0.9128999710083008	
	loss:	0.2987935543060303	

Model: sequential_1-8-1

Loss Function: poisson

	accuracy:	0.9125999808311462	
	loss:	0.1306813806295395	

Model: sequential_1-8-1

Loss Function: binary_crossentropy

	accuracy:	0.902899980545044	
	loss:	0.07139772921800613	

Model: sequential_1-8-2

Loss Function: categorical_crossentropy

	accuracy:	0.9584000110626221	
	loss:	0.15162043273448944	

Model: sequential_1-8-2

Loss Function: poisson

	accuracy:	0.9440000057220459	
	loss:	0.11986614763736725	

Model: sequential_1-8-2
 Loss Function: binary_crossentropy
 | accuracy: 0.9598000049591064 |
 | loss: 0.03655649349093437 |

Model: sequential_1-8-3
 Loss Function: categorical_crossentropy
 | accuracy: 0.9678000211715698 |
 | loss: 0.10980123281478882 |

Model: sequential_1-8-3
 Loss Function: poisson
 | accuracy: 0.9702000021934509 |
 | loss: 0.11022289842367172 |

Model: sequential_1-8-3
 Loss Function: binary_crossentropy
 | accuracy: 0.9629999995231628 |
 | loss: 0.03193367272615433 |

Model: sequential_1-16-1
 Loss Function: categorical_crossentropy
 | accuracy: 0.9176999926567078 |
 | loss: 0.2894260883331299 |

Model: sequential_1-16-1
 Loss Function: poisson
 | accuracy: 0.9151999950408936 |
 | loss: 0.12938936054706573 |

Model: sequential_1-16-1
 Loss Function: binary_crossentropy
 | accuracy: 0.9031000137329102 |
 | loss: 0.06957709044218063 |

Model: sequential_1-16-2
 Loss Function: categorical_crossentropy
 | accuracy: 0.9653000235557556 |
 | loss: 0.12096966803073883 |

Model: sequential_1-16-2
 Loss Function: poisson

	accuracy:	0.9696000218391418	
	loss:	0.11008027195930481	

Model: sequential_1-16-2
 Loss Function: binary_crossentropy

	accuracy:	0.9588000178337097	
	loss:	0.03415116295218468	

Model: sequential_1-16-3
 Loss Function: categorical_crossentropy

	accuracy:	0.979200005531311	
	loss:	0.06647966057062149	

Model: sequential_1-16-3
 Loss Function: poisson

	accuracy:	0.9785000085830688	
	loss:	0.10694978386163712	

Model: sequential_1-16-3
 Loss Function: binary_crossentropy

	accuracy:	0.9753999710083008	
	loss:	0.02192365936934948	

Model: sequential_1-32-1
 Loss Function: categorical_crossentropy

	accuracy:	0.9179999828338623	
	loss:	0.2829427421092987	

Model: sequential_1-32-1
 Loss Function: poisson

	accuracy:	0.9153000116348267	
	loss:	0.12937746942043304	

Model: sequential_1-32-1
 Loss Function: binary_crossentropy

	accuracy:	0.9060999751091003	
	loss:	0.06932933628559113	

Model: sequential_1-32-2
 Loss Function: categorical_crossentropy
 | accuracy: 0.9764000177383423 |
 | loss: 0.07710807770490646 |

Model: sequential_1-32-2
 Loss Function: poisson
 | accuracy: 0.9769999980926514 |
 | loss: 0.10761921107769012 |

Model: sequential_1-32-2
 Loss Function: binary_crossentropy
 | accuracy: 0.9721999764442444 |
 | loss: 0.023090384900569916 |

Model: sequential_1-32-3
 Loss Function: categorical_crossentropy
 | accuracy: 0.9825999736785889 |
 | loss: 0.0563504621386528 |

Model: sequential_1-32-3
 Loss Function: poisson
 | accuracy: 0.982699990272522 |
 | loss: 0.10552532970905304 |

Model: sequential_1-32-3
 Loss Function: binary_crossentropy
 | accuracy: 0.979200005531311 |
 | loss: 0.018165795132517815 |

Model: sequential_2-8-1
 Loss Function: categorical_crossentropy
 | accuracy: 0.8482000231742859 |
 | loss: 0.5114039778709412 |

Model: sequential_2-8-1
 Loss Function: poisson
 | accuracy: 0.8572999835014343 |
 | loss: 0.14986096322536469 |

Model: sequential_2-8-1
 Loss Function: binary_crossentropy

	accuracy:	0.8082000017166138	
	loss:	0.12720641493797302	

Model: sequential_2-8-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9775000214576721	
	loss:	0.07472839951515198	

Model: sequential_2-8-2
 Loss Function: poisson

	accuracy:	0.9786999821662903	
	loss:	0.10688614845275879	

Model: sequential_2-8-2
 Loss Function: binary_crossentropy

	accuracy:	0.975600004196167	
	loss:	0.022512884810566902	

Model: sequential_2-8-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9839000105857849	
	loss:	0.052849579602479935	

Model: sequential_2-8-3
 Loss Function: poisson

	accuracy:	0.9821000099182129	
	loss:	0.1056961640715599	

Model: sequential_2-8-3
 Loss Function: binary_crossentropy

	accuracy:	0.9814000129699707	
	loss:	0.016296260058879852	

Model: sequential_2-16-1
 Loss Function: categorical_crossentropy

	accuracy:	0.8700000047683716	
	loss:	0.4206395149230957	

Model: sequential_2-16-1
 Loss Function: poisson

	accuracy:	0.8639000058174133	
	loss:	0.14427681267261505	

Model: sequential_2-16-1
 Loss Function: binary_crossentropy

	accuracy:	0.857200026512146	
	loss:	0.09680618345737457	

Model: sequential_2-16-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9843000173568726	
	loss:	0.05187729001045227	

Model: sequential_2-16-2
 Loss Function: poisson

	accuracy:	0.9837999939918518	
	loss:	0.10517852753400803	

Model: sequential_2-16-2
 Loss Function: binary_crossentropy

	accuracy:	0.9828000068664551	
	loss:	0.016292624175548553	

Model: sequential_2-16-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9884999990463257	
	loss:	0.036053456366062164	

Model: sequential_2-16-3
 Loss Function: poisson

	accuracy:	0.9876000285148621	
	loss:	0.10367996990680695	

Model: sequential_2-16-3
 Loss Function: binary_crossentropy

	accuracy:	0.9861000180244446	
	loss:	0.010561351664364338	

Model: sequential_2-32-1
 Loss Function: categorical_crossentropy

	accuracy:	0.8700000047683716	
	loss:	0.4233608841896057	

Model: sequential_2-32-1
 Loss Function: poisson

	accuracy:	0.8702999949455261	
	loss:	0.14102916419506073	

Model: sequential_2-32-1
 Loss Function: binary_crossentropy

	accuracy:	0.885200023651123	
	loss:	0.07814329117536545	

Model: sequential_2-32-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9884999990463257	
	loss:	0.03674345090985298	

Model: sequential_2-32-2
 Loss Function: poisson

	accuracy:	0.9865999817848206	
	loss:	0.10388823598623276	

Model: sequential_2-32-2
 Loss Function: binary_crossentropy

	accuracy:	0.9864000082015991	
	loss:	0.012070694006979465	

Model: sequential_2-32-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9876000285148621	
	loss:	0.03534656763076782	

Model: sequential_2-32-3
 Loss Function: poisson

	accuracy:	0.9911999702453613	
	loss:	0.10285915434360504	

Model: sequential_2-32-3
 Loss Function: binary_crossentropy

	accuracy:	0.9901999831199646	
	loss:	0.007893931120634079	

Model: sequential_3-8-1
 Loss Function: categorical_crossentropy

	accuracy:	0.6503999829292297	
	loss:	1.0868866443634033	

Model: sequential_3-8-1
 Loss Function: poisson

	accuracy:	0.6401000022888184	
	loss:	0.21130110323429108	

Model: sequential_3-8-1
 Loss Function: binary_crossentropy

	accuracy:	0.6452000141143799	
	loss:	0.19146348536014557	

Model: sequential_3-8-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9717000126838684	
	loss:	0.0974355936050415	

Model: sequential_3-8-2
 Loss Function: poisson

	accuracy:	0.9739000201225281	
	loss:	0.10887933522462845	

Model: sequential_3-8-2
 Loss Function: binary_crossentropy

	accuracy:	0.9664999842643738	
	loss:	0.02884790301322937	

Model: sequential_3-8-3
 Loss Function: categorical_crossentropy

	accuracy:	0.9714999794960022	
	loss:	0.09493064880371094	

Model: sequential_3-8-3
 Loss Function: poisson

	accuracy:	0.9696999788284302	
	loss:	0.11098623275756836	

Model: sequential_3-8-3
 Loss Function: binary_crossentropy

	accuracy:	0.9646999835968018	
	loss:	0.030567891895771027	

Model: sequential_3-16-1
 Loss Function: categorical_crossentropy

	accuracy:	0.6650999784469604	
	loss:	1.0373120307922363	

Model: sequential_3-16-1
 Loss Function: poisson

	accuracy:	0.6574000120162964	
	loss:	0.20624008774757385	

Model: sequential_3-16-1
 Loss Function: binary_crossentropy

	accuracy:	0.650600016117096	
	loss:	0.1845005750656128	

Model: sequential_3-16-2
 Loss Function: categorical_crossentropy

	accuracy:	0.9825000166893005	
	loss:	0.05720340088009834	

Model: sequential_3-16-2
 Loss Function: poisson

	accuracy:	0.9825999736785889	
	loss:	0.10575384646654129	

Model: sequential_3-16-2
 Loss Function: binary_crossentropy

	accuracy:	0.9799000024795532	
	loss:	0.016392391175031662	

Model: sequential_3-16-3
 Loss Function: categorical_crossentropy
 | accuracy: 0.9807999730110168 |
 | loss: 0.06029291823506355 |

Model: sequential_3-16-3
 Loss Function: poisson
 | accuracy: 0.9812999963760376 |
 | loss: 0.10625935345888138 |

Model: sequential_3-16-3
 Loss Function: binary_crossentropy
 | accuracy: 0.9819999933242798 |
 | loss: 0.014470896683633327 |

Model: sequential_3-32-1
 Loss Function: categorical_crossentropy
 | accuracy: 0.7027000188827515 |
 | loss: 0.9355225563049316 |

Model: sequential_3-32-1
 Loss Function: poisson
 | accuracy: 0.6965000033378601 |
 | loss: 0.19633089005947113 |

Model: sequential_3-32-1
 Loss Function: binary_crossentropy
 | accuracy: 0.666100025177002 |
 | loss: 0.1765071451663971 |

Model: sequential_3-32-2
 Loss Function: categorical_crossentropy
 | accuracy: 0.9879000186920166 |
 | loss: 0.03969316557049751 |

Model: sequential_3-32-2
 Loss Function: poisson
 | accuracy: 0.9889000058174133 |
 | loss: 0.10389383882284164 |

```

Model: sequential_3-32-2
Loss Function: binary_crossentropy
|      accuracy:      0.986299991607666      |
|      loss:          0.011301561258733273    |

```

```

Model: sequential_3-32-3
Loss Function: categorical_crossentropy
|      accuracy:      0.9872000217437744      |
|      loss:          0.04405169561505318      |

```

```

Model: sequential_3-32-3
Loss Function: poisson
|      accuracy:      0.9873999953269958      |
|      loss:          0.10412595421075821      |

```

```

Model: sequential_3-32-3
Loss Function: binary_crossentropy
|      accuracy:      0.9861999750137329      |
|      loss:          0.010200412943959236      |

```

Finally done :)

The following is the best performing model from our tests:

```
[ ]: best_performance(best_model)
```

```
Model: "sequential_2-32-3"
```

Layer (type)	Output Shape	Param #
conv2d_162 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_162 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_163 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_163 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_81 (Flatten)	(None, 1600)	0
dropout_81 (Dropout)	(None, 1600)	0

dense_81 (Dense) (None, 10) 16010

=====

Total params: 34826 (136.04 KB)
Trainable params: 34826 (136.04 KB)
Non-trainable params: 0 (0.00 Byte)

Using binary crossentropy loss function resulted in the following accuracy and loss

	accuracy:	0.9901999831199646	
	loss:	0.007893931120634079	

Layer (type) Output Shape Param #
=====

conv2d_162 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_162 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_163 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_163 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_81 (Flatten)	(None, 1600)	0
dropout_81 (Dropout)	(None, 1600)	0
dense_81 (Dense)	(None, 10)	16010

=====

Total params: 34826 (136.04 KB)
Trainable params: 34826 (136.04 KB)
Non-trainable params: 0 (0.00 Byte)

Using binary crossentropy loss function resulted in the following accuracy and loss

	accuracy:	0.9901999831199646	
	loss:	0.007893931120634079	